

# Assignment 2: Time Series Analysis – Fundamental Concepts

*Andrew G. Dunn*

*January 21, 2016*

**Andrew G. Dunn, Northwestern University Predictive Analytics Program**

Prepared for PREDICT-413: Time Series Analysis and Forecasting.

Formatted using the L<sup>A</sup>T<sub>E</sub>X, via pandoc and R Markdown. References managed using pandoc-citeproc.

## Setup

```
require(fBasics)    # for calculations
require(fpp)        # for data
require(knitr)       # for table output
require(ggplot2)     # for graphing
require(ggthemes)    # for graphing beautifully
require(gridExtra)  # for laying out graphs
```

## Part 1

Consider the monthly returns for General Electric (GE) stock, Center for Research In Security Prices (CRSP) value-weighted index (VW), CRSP equal-weighted index (EW), and S&P composite index (SP) from January 1981 to December 2013. The returns include dividend distributions. Data file is `m-ge3dx8113.txt` with column names `permno` of GE, `date`, `ge`, `vwret`, `ewret`, and `sprtrn`, respectively.

```
d1 = read.table("data/m-ge3dx8113.txt", header=T)
head(d1)
```

```
##   PERMNO    date      ge   vwret   ewret   sprtrn
## 1  12060 19810130 0.000000 -0.040085 0.005615 -0.045742
## 2  12060 19810227 0.089796 0.015521 0.002150 0.013277
## 3  12060 19810331 0.014981 0.046184 0.072674 0.036033
## 4  12060 19810430 -0.020522 -0.011268 0.027885 -0.023456
## 5  12060 19810529 0.001905 0.013551 0.027187 -0.001657
## 6  12060 19810630 -0.046768 -0.010242 -0.013194 -0.010408
```

## Part A

Compute the sample mean, standard deviation, skewness, excess kurtosis, minimum, and maximum of each simple return series.

```
d1a_stats = basicStats(d1)
kable(d1a_stats[c('Mean', 'Stdev', 'Skewness', 'Kurtosis', 'Minimum', 'Maximum'), -(1:2)],
      caption='Basic Statistics of the Simple Return Series')
```

Table 1: Basic Statistics of the Simple Return Series

	ge	vwret	ewret	sprtrn
Mean	0.012900	0.009698	0.011022	0.007594
Stdev	0.071073	0.045036	0.053461	0.043921
Skewness	-0.226160	-0.780736	-0.499120	-0.658830
Kurtosis	1.373376	2.526277	3.259182	2.204877
Minimum	-0.272877	-0.225363	-0.272248	-0.217630
Maximum	0.251236	0.128496	0.225012	0.131767

## Part B

Transform the simple returns to log returns. Compute the sample mean, standard deviation, skewness, excess kurtosis, minimum, and maximum of each log return series.

```
d1b = log(d1[,-(1:2)]+1) # Log Transform, +1 as an offset so that we don't compute log(0)
d1b_stats=ts = basicStats(d1b)
kable(d1b_stats[c('Mean', 'Stdev', 'Skewness', 'Kurtosis', 'Minimum', 'Maximum'),],
      caption='Basic Statistics of the Log Transformed Simple Return Series')
```

Table 2: Basic Statistics of the Log Transformed Simple Return Series

	ge	vwretd	ewretd	sprtrn
Mean	0.010307	0.008630	0.009529	0.006594
Stdev	0.071445	0.045617	0.054030	0.044422
Skewness	-0.615128	-1.076565	-0.942621	-0.933307
Kurtosis	2.273898	3.817684	4.784350	3.276645
Minimum	-0.318660	-0.255361	-0.317795	-0.245428
Maximum	0.224132	0.120886	0.202951	0.123780

## Part C

Test the null hypothesis that the mean of the log returns of GE stock is zero.

```
t.test(d1b$ge)

##
## One Sample t-test
##
## data: d1b$ge
## t = 2.8708, df = 395, p-value = 0.004315
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
## 0.003248467 0.017365132
## sample estimates:
## mean of x
## 0.0103068
```

Reject the null hypothesis that the mean of the log return of GE stock is zero at the 0.05 level, based on p-value.

## Part D

Obtain the empirical density plot of the daily log returns of GE stock and the S&P composite index.

```
pge = ggplot(d1b, aes(ge)) +
  stat_density(alpha = 0.4) +
  labs(x="Returns", y="Density") +
```

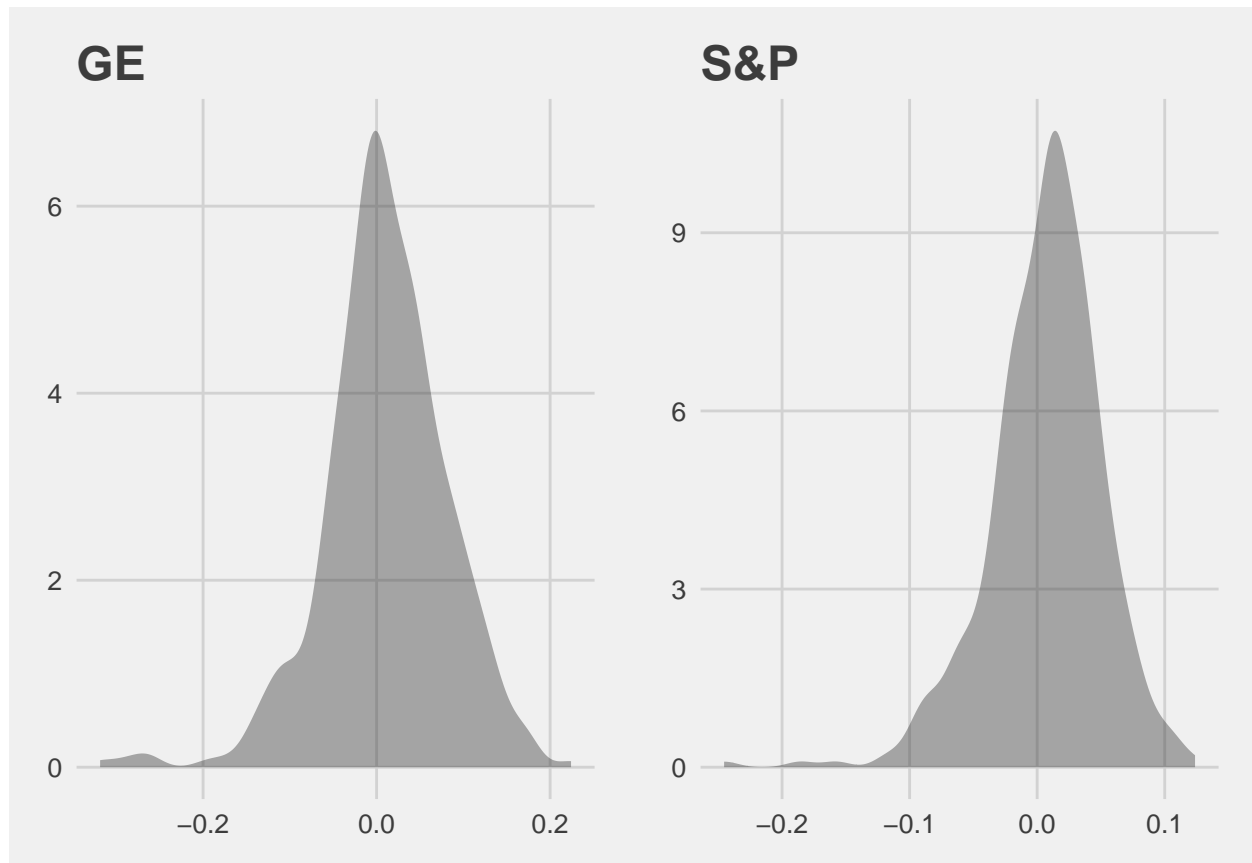
```

ggtitle("GE") + theme_fivethirtyeight()

psprtrn = ggplot(d1b, aes(sprtrn)) +
  stat_density(alpha = 0.4) +
  labs(x="Returns", y="Density") +
  ggtitle("S&P") + theme_fivethirtyeight()

grid.arrange(pge, psprtrn, ncol=2)

```



## Part 2

Consider the daily log returns of Netflix stock from January 2, 2009 to December 31, 2013 as in Problem 1, Assignment 1. Perform the following tests:

```
d2 = read.table("data/d-nflx3dx0913.txt", header=T)
head(d2)
```

```
##   PERMNO    date      nflx    vwretd    ewretd    sprtrn
## 1  89393 20090102 -0.000669  0.030501  0.038274  0.031608
## 2  89393 20090105  0.069300 -0.000580  0.016764 -0.004668
## 3  89393 20090106  0.031309  0.011297  0.033647  0.007817
## 4  89393 20090107 -0.006982 -0.030489 -0.022271 -0.030010
## 5  89393 20090108  0.013452  0.006283  0.011896  0.003397
## 6  89393 20090109 -0.026848 -0.022410 -0.018748 -0.021303
```

## Part A

Test the null hypothesis that the log return is symmetric with respect to its mean;

```
d2l = log(d2[,-(1:2)] + 1) # Log Transform, +1 as an offset so that we don't compute log(0)
st = skewness(d2l$nflx) / sqrt(6 / length(d2l$nflx)) # compute skewness test
print(paste("Skewness Statistic: ", st))
```

```
## [1] "Skewness Statistic: -6.29427781659625"
```

```
p_st = 2 * (1 - pnorm(abs(st))) # computing the p-value
print(paste("p-value: ", p_st))
```

```
## [1] "p-value: 3.08834291473659e-10"
```

Test  $H_0 : M_3 = 0$  versus  $H_a : M_3 \neq 0$ , where  $M_3$  denotes the skewness of the return. Reject the null hypothesis at a 0.05 level based on p-value.

## Part B

Test the null hypothesis that the excess kurtosis of the returns is zero;

Test  $H_0 : K = 3$  versus  $H_a : K \neq 3$ , where  $K$  denotes the kurtosis.

```
kt = kurtosis(d2l$nflx) / sqrt(24 / length(d2l$nflx)) # compute kurtosis test
print(paste("Kurtosis Statistic: ", kt))
```

```
## [1] "Kurtosis Statistic: 170.499517930473"
```

```
p_kt = 2 * (1 - pnorm(abs(kt)))
print(paste("p-value: ", p_kt))
```

```
## [1] "p-value: 0"
```

Reject null hypothesis at a 0.05 level.

## Part C

Construct a 95% confidence interval for the expected daily log return of Netflix stock.

```
t_test = t.test(d2l$nlx)
print(paste("Confidence Interval, Confidence Level 95%: ", t_test[4]))
```

```
## [1] "Confidence Interval, Confidence Level 95%: c(-0.000126484354881271, 0.00411859146786966)"
```

## Part 3


For this exercise, use the quarterly UK passenger vehicle production data from 1997:1 to 2005:1 (data set ukcars) from the Hyndeman text.

```
d3 = window(ukcars, start=1997)
```

### Part A

Plot the data and describe the main features of the series.

```
plot(d3, type="o", xlab = "Years", ylab = "UK Car Production (In Thousands)")
```




The plot indicates the presence of seasonality. It appears that Q3 would be the yearly low, aside from 2002 where the fourth quarter is lower than the third.

### Part B

Decompose the series using STL and obtain the seasonally adjusted data.

```
fit_stl = stl(d3, s.window = "periodic")  
plot(fit_stl)
```



The plot indicates that the seasonal fluctuations do not vary with the level of the time series. The smoothed trend plot indicates a gradual up-trend until just before 2000, with a steep down-trend starting around 2000, then a gradual up-trend starting around 2001.

```
seas_adj = seasadj(fit_stl)
seas_factors = fit_stl$time.series[2:11, "seasonal"] # Acquire the seasonal Factors
```



## Part C


Forecast the next two years of the series using an additive damped trend method applied to the seasonally adjusted data. Then reseasonalize the forecasts. Record the parameters of the method and report the RMSE of the one-step forecasts from your method.

```
fit_damped_seas_adj = holt(seas_adj, damped = TRUE)
print(fit_damped_seas_adj)
```

##	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
##	2005 Q2	402.7663	373.6921	431.8404	358.3011	447.2314
##	2005 Q3	402.6509	365.3414	439.9604	345.5910	459.7109
##	2005 Q4	402.5587	358.3842	446.7332	334.9996	470.1177
##	2006 Q1	402.4849	352.2749	452.6948	325.6954	479.2743
##	2006 Q2	402.4258	346.7575	458.0942	317.2884	487.5632
##	2006 Q3	402.3786	341.6862	463.0709	309.5576	495.1995
##	2006 Q4	402.3408	336.9687	467.7129	302.3628	502.3188
##	2007 Q1	402.3105	332.5411	472.0800	295.6074	509.0137
##	2007 Q2	402.2864	328.3574	476.2153	289.2218	515.3509
##	2007 Q3	402.2670	324.3827	480.1514	283.1532	521.3808

```
plot(fit_damped_seas_adj, xlab = "Years", ylab = "UK Car Production (In Thousands)")
```

### Forecasts from Damped Holt's method



```
print(fit_damped_seas_adj$model)
```


```
## Damped Holt's method
##
## Call:
## holt(x = seas_adj, damped = TRUE)
##
## Smoothing parameters:
##   alpha = 0.7781
##   beta  = 0.0181
##   phi   = 0.8
##
## Initial states:
##   l = 409.7158
##   b = 10.8562
##
## sigma: 22.6867
##
##      AIC      AICc      BIC
## 331.4222 333.6444 338.9047
```

```
print(accuracy(fit_damped_seas_adj))
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -1.815098 22.6867 17.97465 -0.6075613 4.365406 0.5728646
##              ACF1
## Training set -0.02372109
```

```
resea_fit_damed_seas_adj = fit_damped_seas_adj$mean + seas_factors # reseasonalize the forecasted data
```

```
plot(d3, type = "o", xlab = "Years", ylab = "UK Car Production (In Thousands)", xlim = c(1997, 2008))
lines(resea_fit_damed_seas_adj, type = "o", col = "blue")
```



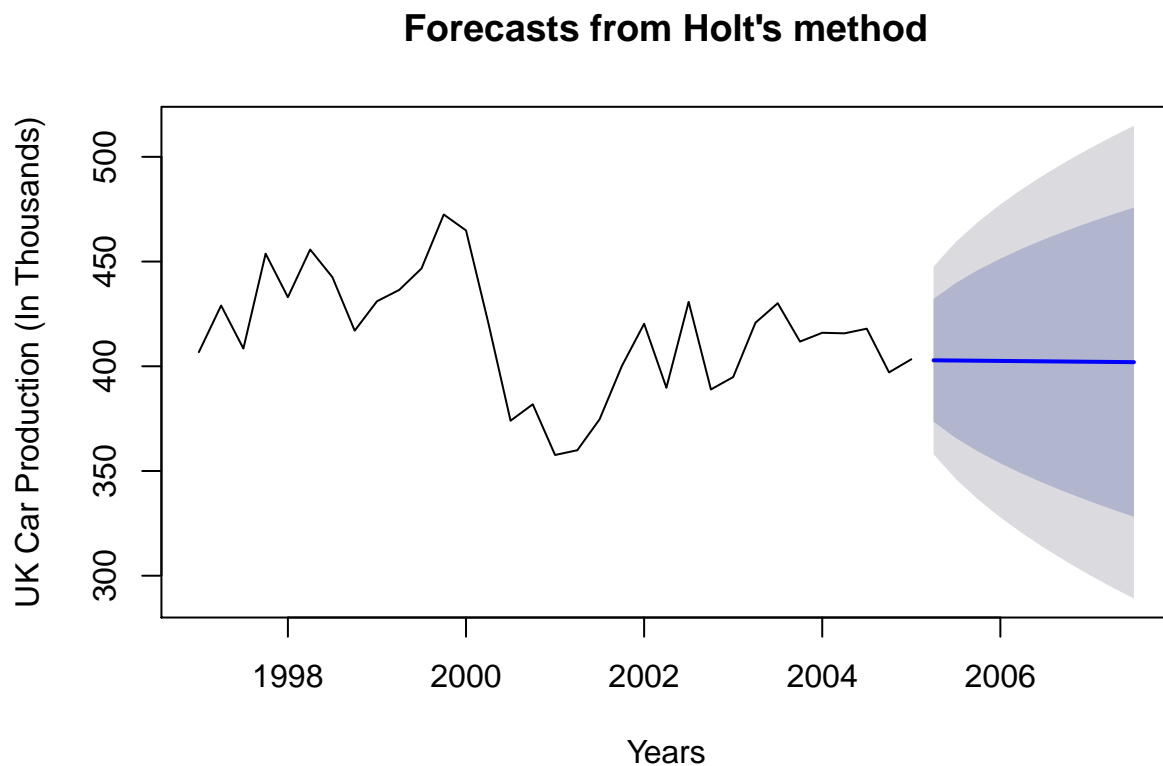
## Part D

Forecast the next two years of the series using Holt's linear method applied to the seasonally adjusted data. Then reseasonalize the forecasts. Record the parameters of the method and report the RMSE of the one-step forecasts from your method.

```
fit_linear = holt(seas_adj)
print(fit_linear)
```

##	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
##	2005 Q2	402.8474	373.5357	432.1591	358.0190	447.6758
##	2005 Q3	402.7490	365.7538	439.7443	346.1697	459.3283
##	2005 Q4	402.6507	359.3119	445.9895	336.3698	468.9316
##	2006 Q1	402.5523	353.6853	451.4193	327.8167	477.2880
##	2006 Q2	402.4540	348.6223	456.2857	320.1255	484.7825
##	2006 Q3	402.3556	343.9788	460.7324	313.0760	491.6353
##	2006 Q4	402.2573	339.6634	464.8511	306.5283	497.9862
##	2007 Q1	402.1589	335.6138	468.7040	300.3869	503.9309
##	2007 Q2	402.0605	331.7850	472.3361	294.5834	509.5377
##	2007 Q3	401.9622	328.1436	475.7807	289.0665	514.8579

```
plot(fit_linear, xlab = "Years", ylab = "UK Car Production (In Thousands)")
```




```
print(fit_linear$model)
```

```
## Holt's method
##
## Call:
## holt(x = seas_adj)
##
## Smoothing parameters:
##   alpha = 0.7698
##   beta  = 1e-04
##
## Initial states:
##   l = 426.7218
##   b = -0.0957
##
## sigma: 22.872
##
##      AIC      AICc      BIC
## 329.9591 331.3877 335.9452
```

```
print(accuracy(fit_linear))
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.8100727 22.87203 18.18449 -0.3841676 4.412324 0.5795523
##              ACF1
## Training set -0.007837827
```

```
resea_linear = fit_linear$mean + seas_factors
plot(d3, type = "o", xlab = "Years", ylab = "UK Car Production (In Thousands)", xlim = c(1997, 2008))
lines(resea_linear, type = "o", col = "blue")
```



## Part E


Now use ETS to choose a seasonal model for the data.

```
fit_ets = ets(d3, model = "ZZZ")
print(fit_ets)
```

```
## ETS(M,N,M)
##
## Call:
## ets(y = d3, model = "ZZZ")
##
## Smoothing parameters:
##   alpha = 0.777
##   gamma = 1e-04
##
## Initial states:
##   l = 423.437
##   s=0.9959 0.9073 1.0288 1.0681
##
## sigma: 0.0549
##
##      AIC      AICc      BIC
## 333.7918 337.0225 342.7708
```

```
plot(forecast(fit_ets), xlab = "Years", ylab = "UK Car Production (In Thousands)")
```

## Forecasts from ETS(M,N,M)



```
print(accuracy(fit_ets))
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.7222169 22.61134 17.9612 -0.3731362 4.396839 0.5724358
##              ACF1
## Training set -0.03674353
```

## Part F

Compare the RMSE of the fitted model with the RMSE of the model you obtained using an STL decomposition with Holt's method. Which gives the better in-sample fits?

```
print(paste("Additive-Damped Model RMSE: ", accuracy(fit_damped_seas_adj)[2]))
```

```
## [1] "Additive-Damped Model RMSE: 22.6867027922175"
```

```
print(paste("Holt Model RMSE: ", accuracy(fit_linear)[2]))
```

```
## [1] "Holt Model RMSE: 22.8720348802071"
```



```
print(paste("ETS Model RMSE: ", accuracy(fit_ets)[2]))
```

```
## [1] "ETS Model RMSE: 22.6113445439822"
```

The ETS model had the lowest RMSE.

## Part G

Compare the forecasts from the two approaches? Which seems most reasonable?

The ETS model seems to be the most reasonable, showing a continuation of the most recent observed trends. We can see, if we compare side by side that both Holt Models show abrupt discontinuation of the recent observed trends.

```
par(mfrow = c(3,1))
p_damped = plot(forecast(fit_damped_seas_adj), type = "o", xlab = "Years", ylab = "Production")
p_linear = plot(forecast(fit_linear), type = "o", xlab = "Years", ylab = "Production")
p_ets = plot(forecast(fit_ets), type = "o", xlab = "Years", ylab = "Production")
```

